



Coding Considerations

Vivado HLS 2013.3 Version

Objectives

➤ After completing this module, you will be able to:

- List language support in Vivado HLS
- State when a construct is not synthesizable
- List unsupported constructs
- Describe multi-access pointer and restriction imposed on using it
- Describe how structs and pointers are handled in Vivado HLS
- Generate efficient hardware from C sources using stream and shift register classes
- List supported functions in OpenCV library

Outline

- ***Language Support***
- **Pointers**
- **Coding Considerations and IO**
- **Streams and Shift Registers**
- **Libraries Support**
 - FFT & FIR
 - OpenCV
- **Summary**

Comprehensive C Support

➤ A Complete C Validation & Verification Environment

- Vivado HLS supports complete bit-accurate validation of the C model
- Vivado HLS provides a productive C-RTL co-simulation verification solution

➤ Vivado HLS supports C, C++ and SystemC

- Functions can be written in any version of C
- Wide support for coding constructs in all three variants of C
 - It's easier to discuss what's not supported than what is

➤ Modeling with bit-accuracy

- Supports arbitrary precision types for all input languages
- Allowing the exact bit-widths to be modeled and synthesized

➤ Floating point support

- Support for the use of float and double in the code

Comprehensive C Support

➤ Pointers based applications

- Multi-access pointer issues and streams

➤ Streaming and shift registering applications

- Generate efficient hardware from C sources using stream and shift register classes

➤ Support for OpenCV functions

- Enable migration of OpenCV designs into Xilinx FPGA
- Libraries target real-time full HD video processing

C, C++ and SystemC Support

► The vast majority of C, C++ and SystemC is supported

- Provided it is statically defined at compile time
- If it's not defined until run time, it won't be synthesizable

► Any of the three variants of C can be used

- If C is used, Vivado HLS expects the file extensions to be .c
- For C++ and SystemC it expects file extensions .cpp

► Let's look at an example using a FIR

- Single data input x
- Coefficients are stored in a ROM
- A single output: function return

FIR Filter with C

➤ C coding style

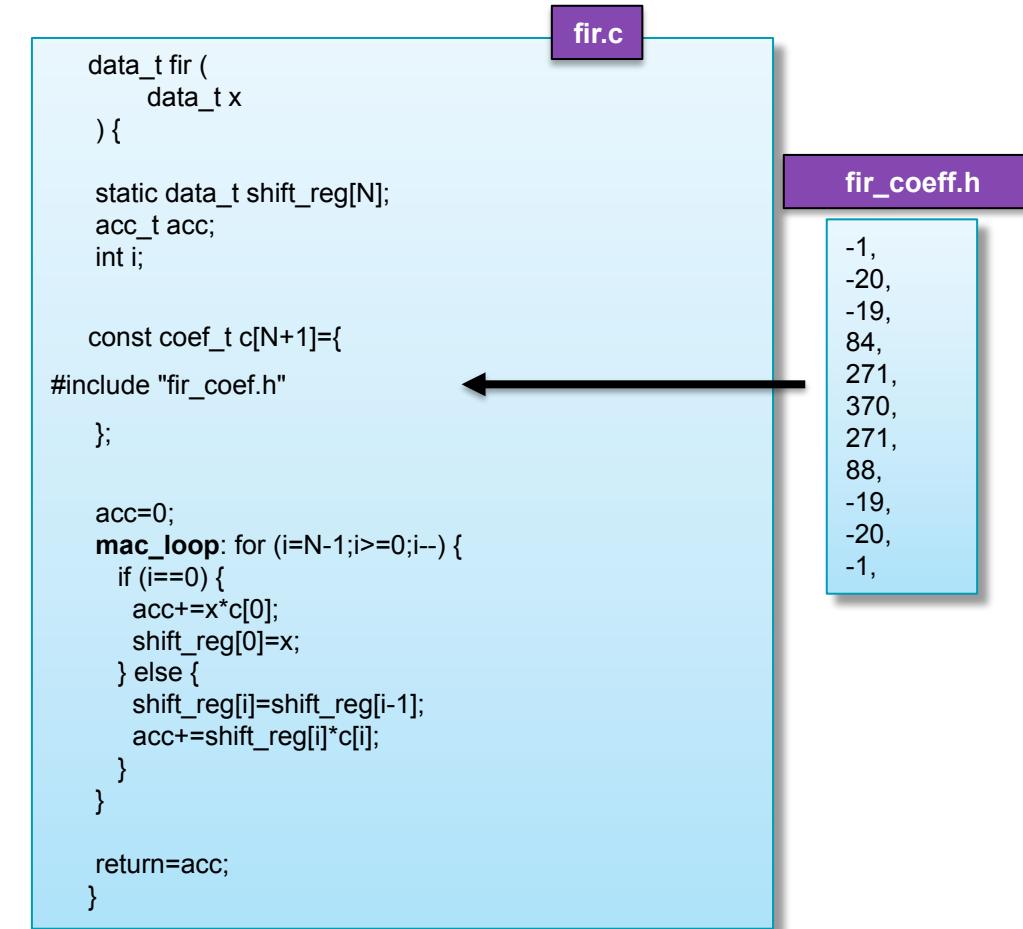
- The top-level function defines the IO ports
- If required, sub-functions can be added

➤ Storage

- Static: required for anything which must hold it's value between invocations
 - A C simulation will confirm what is required to be a static
 - Will become a register or memory
 - Other variables may become storage, depending on the implementation
- Const: ensures the values are seen as constants which never changed

➤ FIR details

- Coefficients are loaded from a file
 - Use the “const” keyword to ensure the data is seen as constant and a ROM is used



Arrays to RAMs : Initialization

➤ Static and Const have great impact on array implementation

- Const, as shown on previous slide, implies a ROM
- Static can impact how RAMs are initialized

➤ A typical coding style changes

```
int coeff[8] = {-2, 8, -4, 10, 14, 10, -4, 8, -2};
```



- Implies coeff is set each time the function is executed
- Resets in the RAM being explicitly loaded each transaction
 - Costs clock cycles
- The array can be initialized as a static

```
static int coeff[8] = {4, -2, 8, -4, 10, 14, 10, -4, 8, -2, 4};
```



- Statics are initialized at “start up” only
 - In the C, in the RTL via bitstream

A coefficient array does not need to be initialized every time but doing that costs “nothing” in software, unlike hardware

FIR Filter with C++

► C++ Coding Style

- Top-level defines RTL IO
- A class defines the functionality
 - Additional classes or sub-functions can be added
- The class is instantiated in the top-level function
 - Methods can be used

```
#include "types.h"
template<class coef_T, class data_T, class acc_T>
class CFir
{
protected:
    static const coef_T c[];
    data_T shift_reg[N-1];
private:
public:
    data_T operator()(data_T x);
    template<class coef_TT, class data_TT, class acc_TT>
};
```

fir.h

CFir Class defined

```
template<class coef_T, class data_T, class acc_T>
const coef_T CFir<coef_T, data_T, acc_T>::c[] =
{
#include "fir_coeff.h"
};

template<class coef_T, class data_T, class acc_T>
data_T CFir<coef_T, data_T, acc_T>::operator()(data_T x)
{
    int i;
    acc_t acc = 0;
    data_t m;
```

CFir functionality defined

```
    mac_loop: for (i=N-1;i>=0;i--) {
        if (i==0) {
            acc+=x*c[0];
            shift_reg[0]=x;
        } else {
            shift_reg[i]=shift_reg[i-1];
            acc+=shift_reg[i]*c[i];
        }
    }
    return acc;
}

data_t fir(data_t x)
{
    static CFir<coef_t, data_t, acc_t> fir1;

    return fir1(x);
}
```

fir.cpp

fir_coeff.h

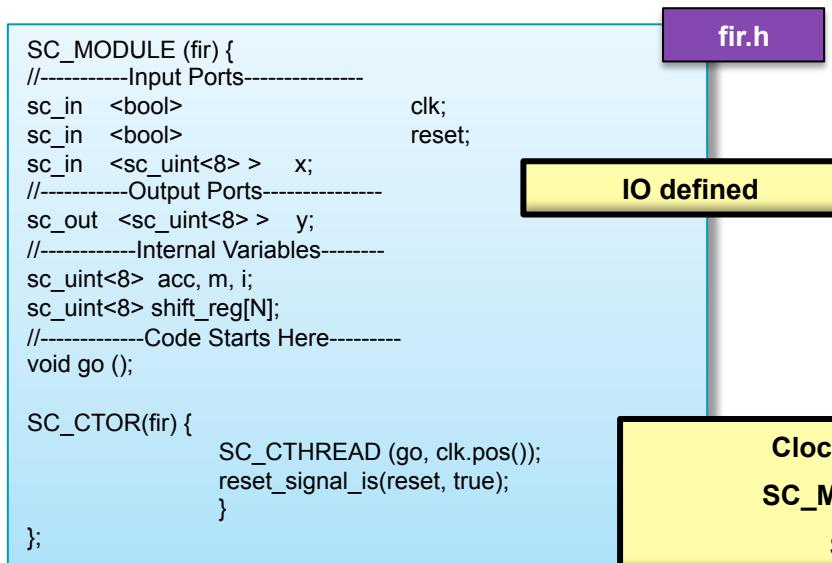
```
-1,
-20,
-19,
84,
271,
370,
271,
88,
-19,
-20,
-1,
```

Top-level function defines the IO Class CFIR is instantiated as fir1

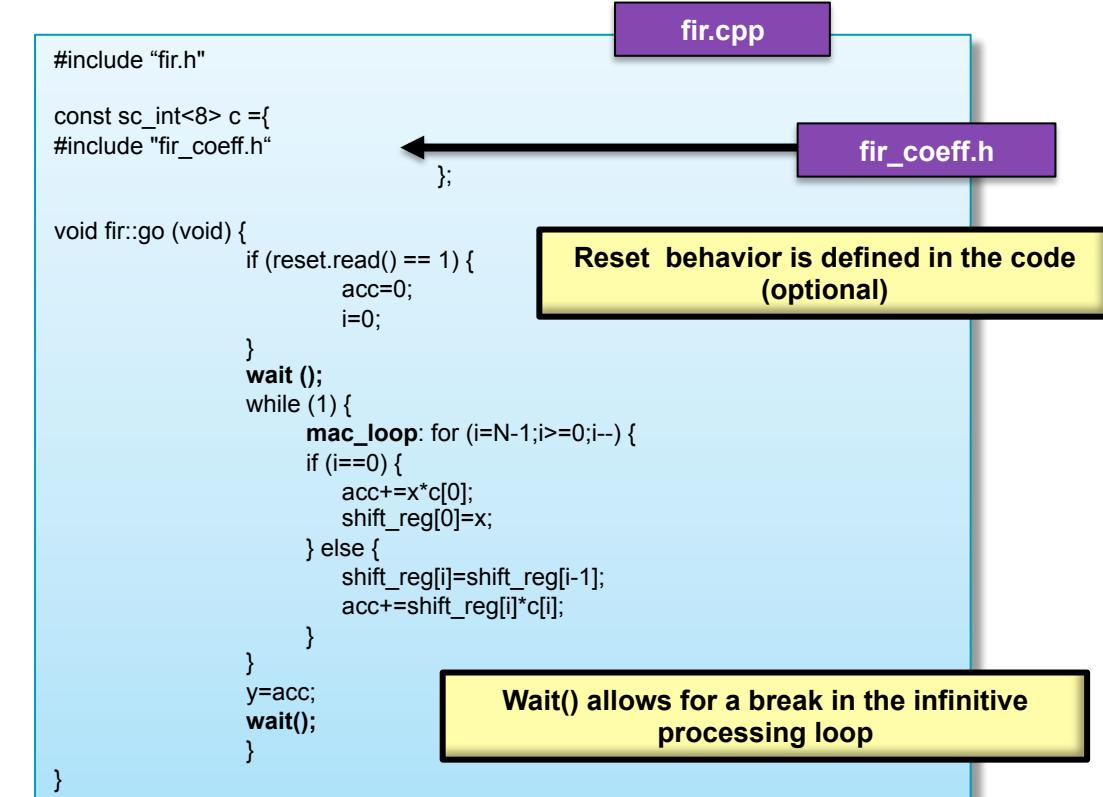
FIR Filter with SystemC

SystemC Style

- The header defines the IO ports
- The constructor calls the functions
 - Can be SC_CTHREAD sensitive on the clock (with optional reset)
 - Can be SC_METHOD sensitive on any signal
- Multiple threads, methods and functions can be used



Clock sensitive thread (SC_CTHREAD)
SC_METHOD is also supported (combo)
SC_THREAD is not supported



#include "fir.h"

const sc_int<8> c =
#include "fir_coeff.h"

void fir::go (void) {

if (reset.read() == 1) {
acc=0;
i=0;

}

wait();
while (1) {

mac_loop: for (i=N-1;i>=0;i--) {
if (i==0) {
acc+=x*c[0];
shift_reg[0]=x;

} else {
shift_reg[i]=shift_reg[i-1];
acc+=shift_reg[i]*c[i];

}

}

y=acc;

wait();

}

fir.cpp

fir_coeff.h

Reset behavior is defined in the code (optional)

Wait() allows for a break in the infinitive processing loop

Unsupported Constructs: Overview

➤ System calls and function pointers

- Dynamic memory allocation
 - malloc() & free()
- Standard I/O and file I/O operations
 - fprintf() / fscanf() etc.
- System calls
 - time(), sleep() etc.

➤ Data types

- Forward declared type
- Recursive type definitions
 - Type contains members with the same type

➤ Non-standard Pointers

- Pointer casting between general data types
 - OK with native integers types
- If a double pointer is used in multiple functions, Vivado HLS will inline all the functions
 - Slower synthesis, may increase area & run time

Unsupported: Dynamic Memory Allocation

► Dynamic memory allocation

- Not allowed since it requires the construction (or destruction) of hardware at runtime
 - *malloc, alloc, free* are **not** synthesizable
 - Similarly for constructor initialization
-
- Use persistent static variables and fixed-size arrays instead

```
long long x = malloc (sizeof(long long));  
int* arr = malloc (64 * sizeof(int));
```

```
static long long x;  
int array[64];
```

► Dynamic virtual function call

```
Class A {  
public:  
    virtual void bar() {...};  
};
```

```
void fun(A* a) {  
    a->bar();  
}  
A* a = 0;  
if (base) A= new A();  
else A = new B();  
foo(a);
```

Unsupported: System Calls

► System calls

- Most of C system calls do not have hardware counterparts and are not synthesizable
 - printf(), getc(), time(), ...
- Vivado HLS will ignore system calls
 - They can also be removed by the `_SYNTHESIS_` macro
- Good Practice:
 - Use Macro “`_SYNTHESIS_`” to remove code segment involving system calls from synthesis
 - Macro “`_SYNTHESIS_`” is automatically defined by the Vivado HLS during elaboration

Only read this code if macro
`_SYNTHESIS_` is not set

```
void foo (...) {  
    ...  
#ifndef _SYNTHESIS_  
    Code will be seen by simulation.  
    But not synthesis.  
#endif  
    ...  
}
```

Unsupported: General Pointer Casting

➤ Pointer reinterpretation

- Casting a pointer to a different type is not allowed in the general case

```
struct {  
    short first;  
    short second;  
} pair;  
  
*(unsigned*)pair = -1U;
```

- Solution: assign values using the original type

```
struct {  
    short first;  
    short second;  
} pair;  
  
pair.first = -1U;  
pair.second = -1U;
```

➤ Pointer casting is allowed between native C integer types

```
int foo (int index, int A[N]) {  
    char* ptr;  
    int i =0, result = 0;  
  
    ptr = (char*)(&A[index]);  
  
    for (i = 0; i < 4*N; ++i) {  
        result += *ptr;  
        ptr+=1;  
    }  
    return result;  
}
```

Unsupported: Recursive Functions

➤ Avoid recursive functions

- Not synthesizable in general
 - The code re-entrance indirectly uses dynamic memory allocation

```
unsigned foo (unsigned n)
{
    if (n == 0 || n == 1) return 1;
    return (foo(n-2) + foo(n-1));
}
```

Not synthesizable: Endless recursion

➤ Standard Template Libraries (STLs)

- Many are unbounded or use recursive functions
 - In general they cannot be synthesized
- Re-code to create a local bounded model
 - Use arrays instead of vectors
 - Shifts instead of queues

Test Bench: Gotcha 1

➤ Use a RETURN value

- RTL verification re-uses the C test bench to verify the RTL design
- If the test bench does not return a 0, RTL sim will say it failed

```
@I [LIC-101] Checked in features [ Vivado HLS_FLOW Vivado HLS_OPT Vivado HLS_SC Vivado HLS_XILINX ]  
Generating autosim.sc.exe  
@I [SIM-6] Starting SystemC simulation ...
```

```
SystemC 2.2.0 --- Sep 15 2011 23:59:06  
Copyright (c) 1996-2006 by all Contributors  
ALL RIGHTS RESERVED
```

Note: VCD trace timescale unit is set by user to 1.000000e-12 sec.

```
0 3  
1 2  
2 1  
3 0
```

```
@E [SIM-3] Simulation failed: test bench return error code "20".  
@E [SIM-1] *** AutoSim finished: FAIL ***
```

The test bench could show the correct results after RTL sim

But if there is no return 0, Vivado HLS will report a failure

➤ Examine the SIM-3 Message

- Message SIM-3 says why the RTL failed verification
- A return of “20” means nothing returned: likely a bad test bench

```
@E [SIM-3] Simulation failed: test bench return error code "20".  
@E [SIM-1] *** AutoSim finished: FAIL ***
```

- A return of “1” shows the self-checking was used: valid failure
 - If the test bench returns “1” when self-checking fails (like previous example)

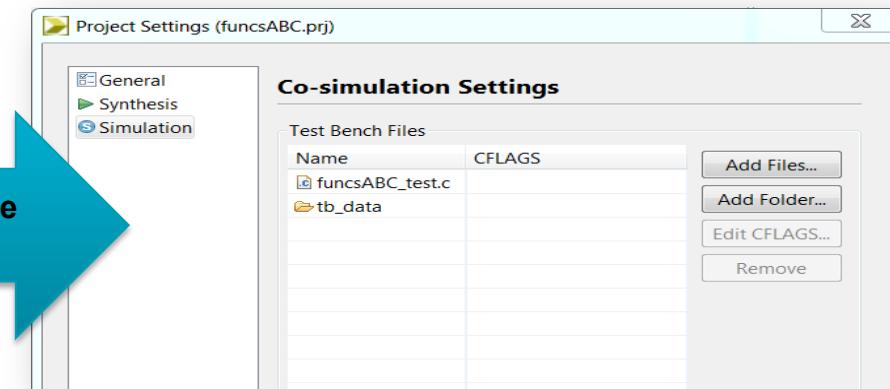
Test Bench : Gotcha 2

➤ Include all files used by the test bench in Vivado HLS

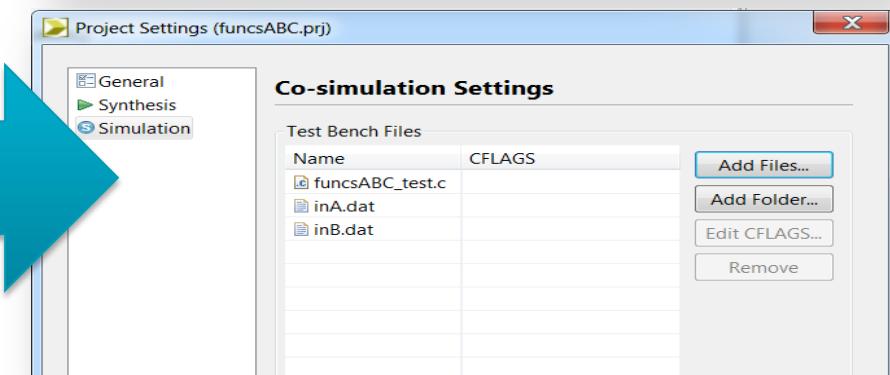
- This test bench reads golden data from files

```
int main () {  
    ...  
    // Get data from files  
    fp=fopen("tb_data/inA.dat","r");  
    for (i=0; i<N*NUM_TRANS; i++){  
        int tmp;  
        fscanf(fp, "%d", &tmp);  
        a_gold[i] = tmp;  
    }  
    fclose(fp);  
  
    fp=fopen("tb_data/inB.dat","r");  
    for (i=0; i<N*NUM_TRANS; i++){  
        int tmp;  
        fscanf(fp, "%d", &tmp);  
        b_gold[i] = tmp;  
    }  
    fclose(fp);  
    ...  
    return ret;  
}
```

Include file directory with the test bench



Or include both files



- Include all files or RTL verification will fail

Outline

- **Language Support**
- **Pointers**
- **Coding Considerations and IO**
- **Streams and Shift Registers**
- **Libraries Support**
 - FFT & FIR
 - OpenCV
- **Summary**

Using Pointer

➤ Structs as pointers

- Structs are implemented differently as pointers or pass-by-value

➤ Pointer arithmetic

- Only supported as interface using ap_bus

➤ Converting pointers using malloc

- Must be converted to fixed sized resources

➤ Multi-access Pointers

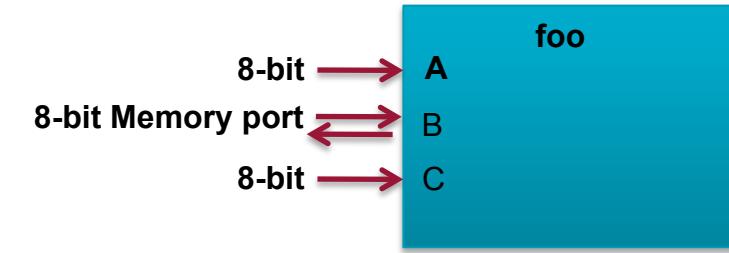
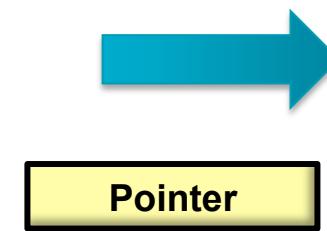
- Must be marked as volatile or reads and writes will be optimized away
- Cannot be rigorously validated with C simulation
- Require a depth specification for RTL simulation
 - These issues were addressed in the previous (IO) section, but re-iterated here

Pointer & Structs

► Struct passed as a pointer

- When a struct is passed as a pointer
- The data will be accessed as pass-by-reference

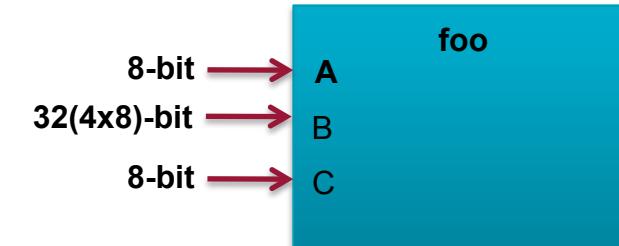
```
typedef struct{  
    unsigned char A;  
    unsigned char B[4];  
    unsigned char C;  
} my_data;  
  
void foo(my_data *in);
```



► Struct passed by-value

- When a struct is passed directly and not as a pointer
- The data will be accessed as pass-by-value

```
typedef struct{  
    unsigned char A;  
    unsigned char B[4];  
    unsigned char C;  
} my_data;  
  
void foo(my_data in);
```



Struct Access

➤ Accessing a struct: pass-by-reference

```
void foo(my_data *a_in, my_data *b_out) {  
    int i;  
  
    b_out->A = a_in->A + 1;  
  
    for(i=0; i <= 319; i++) {  
        b_out->B[i] = a_in->B[i] +1 ;  
    }  
  
    b_out->C = a_in->C +1;  
}
```

```
typedef struct {  
    unsigned char A;  
    unsigned char B[320];  
    unsigned char C;  
}my_data;
```

➤ Accessing a struct: pass-by-value

```
my_data foo(my_data a_in, my_data b_out) {  
    int i;  
  
    b_out.A = a_in.A+1;  
  
    for(i=0; i <= N-1; i++) {  
        b_out.B[i] = a_in.B[i]+1;  
    }  
  
    b_out.C = a_in.C+1;  
  
    return b_out;  
}
```

Pointer Arithmetic: One Interface

➤ Pointer arithmetic

- This function uses some pointer arithmetic

```
#include "bus.h"

void foo (int *d) {
    static int acc = 0;
    int i;

    for (i=0;i<4;i++) {
        acc += *(d+i+1);
        *(d+i) = acc;
    }
}
```

↔
Pointer
arithmetic

Value of “d” are defined in the test
bench

```
# int main () {
    int d[5];
    int i;

    for (i=0;i<5;i++) {
        d[i] = i;
    }

    foo(d);

    // Check results
    ...
}
```

- Argument can be implemented as ap_bus or ap_fifo

➤ With a FIFO interface

- All reads & writes must be sequential (start from location 0 implied)
- This functionality can only be implemented using ap_bus

➤ Address generated by pointer arithmetic

- Can only be exported externally using ap_bus
- All other pointer interfaces are streaming type

Converting Malloc Pointers

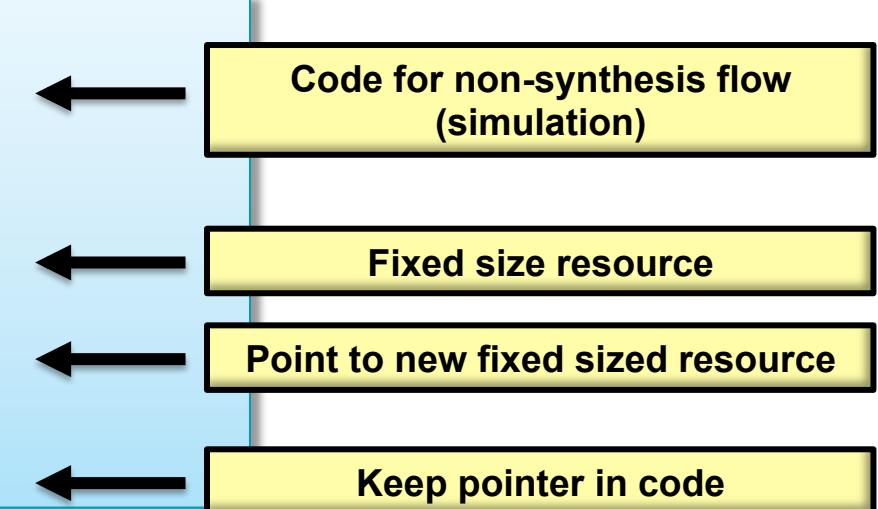
► Many design use Malloc and Pointer

- Malloc is not supported for synthesis
- Must be converted to a resource of a defined size → array
- This can mean changing lots of code: *var → var[]

► The following workaround can ensure limited code impact

- Create a dummy variable
- Set the pointer to the dummy variable

```
// Internal image buffers
#ifndef __SYNTHESIS__
    my_type *yuv = (my_type *)malloc(N*sizeof(my_type));
    my_type *scale = (my_type *)malloc(N*sizeof(my_type));
#else // Workaround malloc() calls w/o changing rest of code
    my_type _yuv[N];
    my_type _scale[N];
    my_type *yuv = &_yuv;
    my_type *scale = &_scale;
#endif
res = *yuv;
```



Outline

- Language Support
- Pointers
- *Coding Considerations and IO*
- Streams and Shift Registers
- Libraries Support
 - FFT & FIR
 - OpenCV
- Summary

Coding and IO

➤ There is one major issue with IO and coding

- Multi-access pointers

➤ A multi-access pointer is a pointer argument which is ...

- Read or written to multiple times

```
void fifo (int *d_o,
           int *d_i
         ) {
    static int acc = 0;
    int cnt;

    for (cnt=0;cnt<4;cnt++) {
        acc += *d_i;
        if (cnt%2==1) {
            *d_o = acc;
        }
    }
}
```

*d_i is read multiple times

*d_o is written to multiple times

- This can result in incorrect hardware
- This cannot be verified using a test bench
- This can give incorrect results at RTL verification

Yes, you can use this coding style (if you
REALLY wish).... With the following
workarounds

Multi-Access Pointers: Use Volatile!!!

➤ Without Volatile

```
void fifo (int *d_o,  
          int *d_i) {  
    static int acc = 0;  
    int cnt;  
    for (cnt=0;cnt<4;cnt++) {  
        acc += *d_i;  
        if (cnt%2==1) {  
            *d_o = acc;  
        }  
    }  
}
```

- C compilers will optimize to 1 read and 1 write
- Vivado HLS will create a design with 1 read and 1 write

To complicate matters, if the IF condition is complex, Vivado HLS may not be able to optimize the accesses into 1

➤ With Volatile

```
void fifo (volatile int *d_o,  
          volatile int *d_i) {  
    static int acc = 0;  
    int cnt;  
    for (cnt=0;cnt<4;cnt++) {  
        acc += *d_i;  
        if (cnt%2==1) {  
            *d_o = acc;  
        }  
    }  
}
```

- C Compilers will not optimize the IO accesses
 - Will assume the pointer value may change outside the scope of the function & leave the accesses
- Vivado HLS will create a design with 4 reads and 2 write

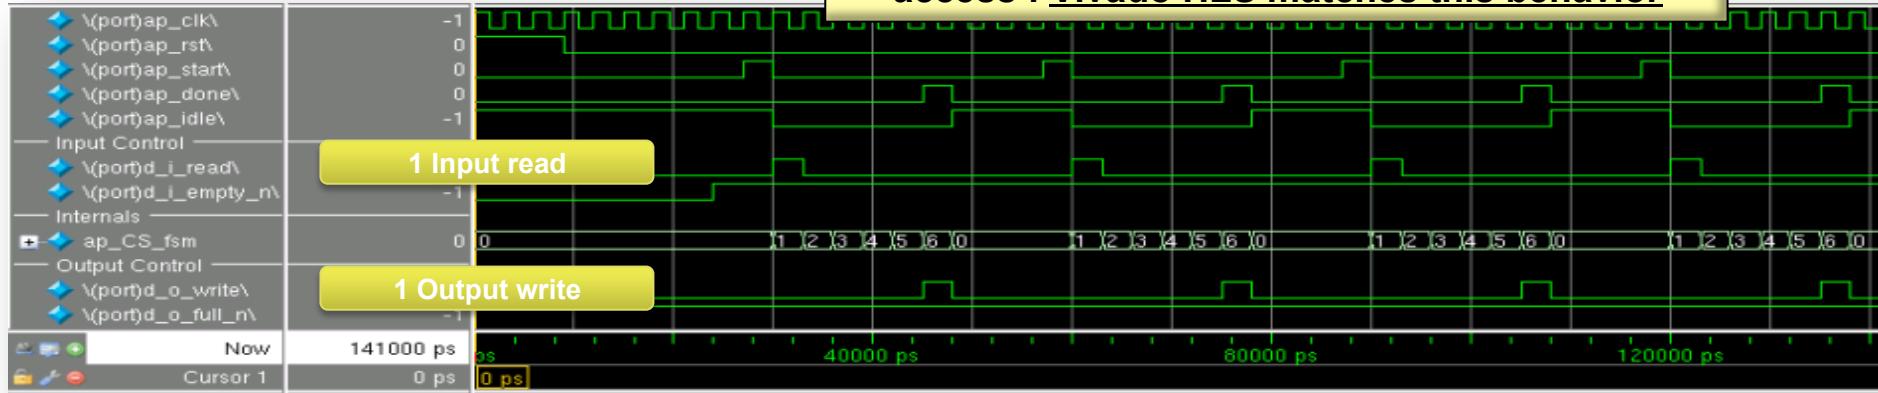
Multi-Access Pointers: the effect of volatile

► Easier to see with a simple example

```
void fifo (
    int *d_o,
    int *d_i
){
    static int acc = 0;
    int cnt;

    acc += *d_i;
    acc += *d_i;
    *d_o = acc;

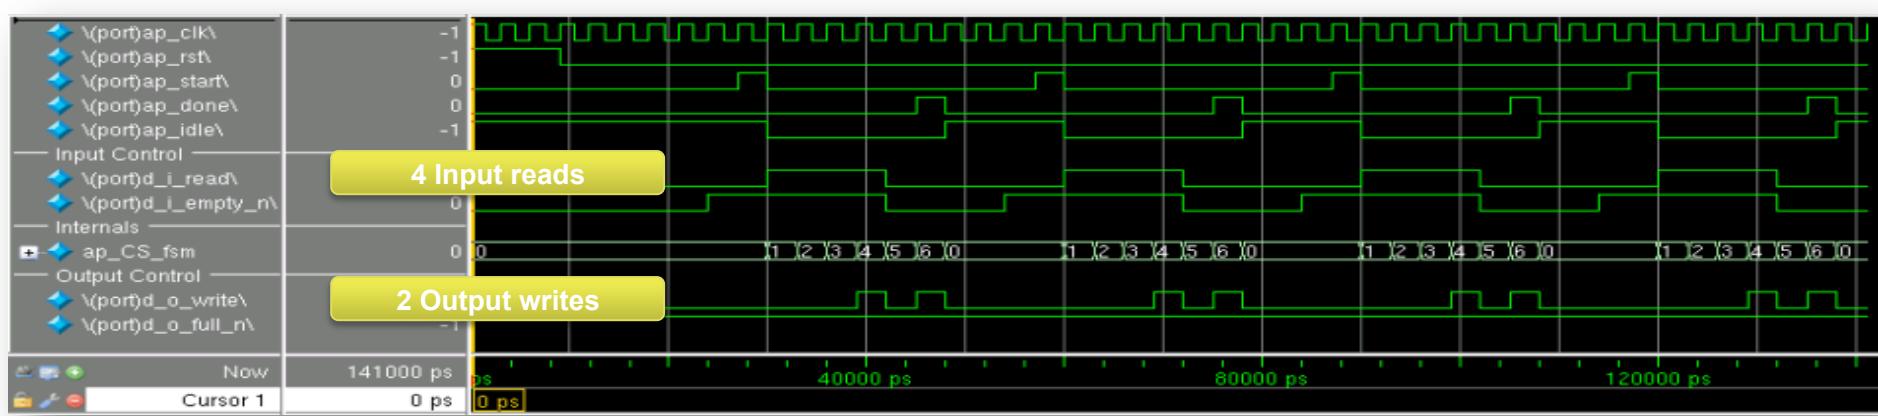
    acc += *d_i;
    acc += *d_i;
    *d_o = acc;
}
```



```
void fifo (
    volatile int *d_o,
    volatile int *d_i
){
    static int acc = 0;
    int cnt;

    acc += *d_i;
    acc += *d_i;
    *d_o = acc;

    acc += *d_i;
    acc += *d_i;
    *d_o = acc;
}
```



C compilers will optimize multiple accesses on the same (non-volatile) pointer into a single access : Vivado HLS matches this behavior

Multi-Access Pointers: Limited Visibility

► The intermediate results are hard to use for verification

- Only the final pointer value is passed to the test bench

```
void fifo (int *d_o,  
          int *d_i  
        ) {  
    static int acc = 0;  
    int cnt;  
  
    for (cnt=0;cnt<4;cnt++) {  
        acc += *d_i;  
        if (cnt%2==1) {  
            *d_o = acc;  
        }  
    }  
}
```

- » The final value of *d_o can be captured in the test bench and compared
- » The intermediate values can only be seen by using printf : they cannot be automatically verified by the test bench

► Code could be added to the DUT

- Use __SYNTHESIS__ to add unsynthesizable code for checking
- Which will not be there in the RTL

Not the most ideal coding style for verification

- HLS:STREAMS class can be used instead: discussed in next section

Multi-Access Pointers: Test bench Depth

► Let's look at an example

- This code will access four samples using a pointer

```
#include "bus.h"

void foo (int *d) {
    static int acc = 0;
    int i;

    for (i=0;i<4;i++) {
        acc += *(d+i);
        *(d+i) = acc;
    }
}
```

The C test bench may correctly provide 4 values

```
int main () {
    int d[5];
    int i;

    for (i=0;i<4;i++) {
        d[i] = i;
    }

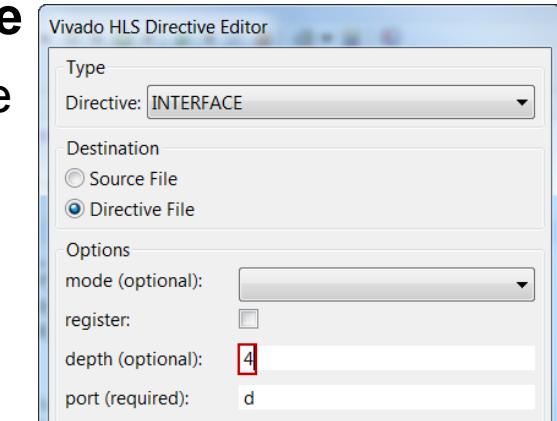
    foo(d);

    return 0;
}
```

► Co-Simulation(RTL Verification) only sees a pointer on the interface

- And will create a SystemC co-simulation wrapper to supply or store 1 sample
- Use the depth option to specify the actual depth
- Vivado HLS will issue a warning in case

@I [SIM-76] 'd' has a depth of '4'. Interface directive option '-depth' can be used to set to a different value. Incorrect depth may result in simulation mismatch.



Outline

- Language Support
- Pointers
- Coding Considerations and IO
- *Streams and Shift Registers*
- Libraries Support
 - FFT & FIR
 - OpenCV
- Summary

Designing with Streams

➤ Streams can be used instead of multi-access pointers

- None of the issues

➤ Streams simulate like an infinite FIFO in software

- Implemented as a FIFO of user-defined size in hardware

➤ Streams have support for multi-access

- Streams interface to the testbench
- Streams can be read in the testbench to check the intermediate values
- Streams store values: no chance of the volatile effect

➤ Streams are supported on the interface and internally

- Define the stream as static to make it internal only

Using Streams

➤ Streams are C++ classes

- Modeled in C++ as an infinite depth FIFO
- Can be written to or read from

➤ Ideal for Hardware Modeling

- Ideal for modeling designs in which the data is known to be streaming (data is always in sequential order)
 - Video or Communication designs typically stream data
- No need to model the design in C++ as a “frame”
- Streams allow it to be modeled as per-sample processing
 - Just fill the stream in the test bench
 - Read and write to the stream as if it's an infinite FIFO
 - Read from the stream in the test bench to confirm results

➤ Stream are by default implemented as a FIFO of depth 2

- Can be specified by the user: needed for decimation/interpolation designs

Stream Example

➤ Create using hls::stream or define the hls namespace

```
#include <ap_int.h>
#include <hls_stream.h>

typedef ap_uint<128> uint128_t;           // 128-bit user defined type

hls::stream<uint128_t> my_wide_stream; // A stream declaration
```

```
#include <ap_int.h>
#include <hls_stream.h>
using namespace hls;                      // Use hls namespace

typedef ap_uint<128> uint128_t;           // 128-bit user defined type

stream<uint128_t> my_wide_stream; // hls:: no longer required
```

➤ Blocking and Non-Block accesses supported

```
// Blocking Write
hls::stream<int> my_stream;

int src_var = 42;
my_stream.write(src_var);
// OR use: my_stream << src_var;
```

```
// Blocking Read
hls::stream<int> my_stream;

int dst_var;
my_stream.read(dst_var);
// OR use: dst_var = my_stream.read();
```

```
hls::stream<int> my_stream;

int src_var = 42;
bool stream_full;

// Non-Blocking Write
if (my_stream.write_nb(src_var)) {
    // Perform standard operations
} else {
    // Write did not happen
}

// Full test
stream_full = my_stream.full();
```

```
hls::stream<int> my_stream;

int dst_var;
bool stream_empty;

// Non-Blocking Read
if (my_stream.read_nb(dst_var)) {
    // Perform standard operations
} else {
    // Read did not happen
}

// Empty test
fifo_empty = my_stream.empty();
```

Stream arrays and structs are not supported for RTL simulation at this time: must be verified manually.

e.g. hls::stream<uint8_t> chan[4])

AP_SHIFT_REG Class

► Class AP_SHIFT_REG infers an SRL

- Ensures that the behavior of an SRL component can be modeled in C
 - Addressable shift register
- Ensures that an SRL is in fact used in the implementation
 - *Guarantees* that the high-performance Xilinx primitive is used

► Using the AP_SHIFT_REG

- For use in C++ designs
- Defined in the Vivado HLS tool header file *ap_shift_reg.h*
 - Must be included

► SRL operations modeled in C

- Supported in the C code
- Addressable read from the shift register
- Addressable read from the shift register and shift
- Addressable read from the shift register and shift if enabled

Using ap_shift_reg class

➤ Read from the shift register

```
// Include the Class  
#include "ap_shift_reg.h"  
  
// Shift reg of 4 integers  
static ap_shift_reg<int, 4> Sreg;  
int var1;  
...  
var1 = Sreg.read(2);  
...
```

Read location 2 of the register into variable “var1”

➤ Read from and write to the shift register

```
// Include the Class  
#include "ap_shift_reg.h"  
  
// Shift reg of 4 integers  
static ap_shift_reg<int, 4> Sreg;  
int var1;  
...  
var1 = Sreg.shift(ln1,2);  
...
```

Read location 2 of the register into variable “var1”

Write variable “in1” into location 0 and shift all values along by 1 (location 3 is lost after this shift)

Using ap_shift_reg class

➤ Read from and optionally write to the shift register

```
// Include the Class  
#include "ap_shift_reg.h"  
  
// Shift reg of 4 integers  
static ap_shift_reg<int, 4> Sreg;  
int var1, In1;  
bool En;  
  
// Read location 3 of Sreg into var1  
// THEN if En=1  
// Shift all values up one and load In1 into location 0  
var1 = Sreg.shift(In1,3,En);
```

Read location 3 of the register into variable “var1”

AND THEN IF signal “En” is active high

Write variable “in1” into location 0 and shift all values along by 1 (location 3 is lost after this shift)

➤ Shift only

```
// Include the Class  
#include "ap_shift_reg.h"  
  
// Shift reg of 4 integers  
static ap_shift_reg<int, 4> Sreg;  
int var1, In1;  
bool En;  
  
Sreg.shift(In1);
```

Write variable “in1” into location 0 and shift all values along by 1 (location 3 is lost after this shift)

Outline

- Language Support
- Pointers
- Coding Considerations and IO
- Streams and Shift Registers
- *Libraries Support*
 - FFT & FIR
 - OpenCV
- Summary

FFT and FIR IP in HLS (Starting in 2013.3)

➤ The Xilinx FFT and FIR IP are available in Vivado HLS

- C simulates with a bit-accurate model
- Fully configurable within the C++ source code
 - Pre-defined C++ structs allow the IP to be configured & accessed

➤ Supported only for C++

- Implemented with templates

➤ High-Quality Implementation

- Same hardware as implemented by RTL versions of this IP
- Functionality fully described in Xilinx Documentation
 - LogiCORE IP Fast Fourier Transform v9.0 (document PG109)
 - LogiCORE IP FIR Compiler v7.1 (document PG149)

FFT Function

➤ Using the FFT

- Include the `hls_fft.h` library in the code
 - This defines the FFT and supporting structs and types
 - Allows `hls::fft` to be instantiated in your code
- Use the `STATIC_PARAM` template parameter to parameterize the FFT
 - The `STATIC_PARAM` template parameter defines all static configuration values
 - The Library provides a pre-defined struct `hls::ip_fft::params_t` to perform this
- Optionally modify the default parameters by creating a new user defined `STATIC_PARAM` struct based on the default

```
#include "hls_fft.h"  
  
hls::fft<STATIC_PARAM> (  
    INPUT_DATA_ARRAY,  
    OUTPUT_DATA_ARRAY,  
    OUTPUT_STATUS,  
    INPUT_RUN_TIME_CONFIGURATION);  
  
// Static Parameterization Struct  
// Input data fixed or float  
// Output data fixed or float  
// Output Status  
// Input Run Time Configuration
```

FFT Static Configuration

► Default Configuration Pre-defined

- The `hls_fft.h` header file defines struct `hls::ip_fft::params_t`
- This contains the default values for configuring the FFT

► User defined struct

- Allows the FFT to be easily configured
- Based on `hls::ip_fft::params_t`
- User struct simply updates any specified value

```
#include "hls_fft.h"

struct param1 : hls::ip_fft::params_t {
    static const unsigned ordering_opt = hls::ip_fft::natural_order;
    static const unsigned config_width = FFT_CONFIG_WIDTH;
    static const unsigned status_width = FFT_STATUS_WIDTH;
};

// FFT IP
hls::fft<param1> (xn1, xk1, &fft_status1, &fft_config1);
```

Include Header File

Create struct based on `hls::ip_fft::params_t`

- `params_t` includes the default values
- Struct `param1` overrides 3 default values in this example
- Struct `param1` is then used to configure the FFT
- Used as `STATIC_PARAM` in previous slide

FFT Configuration Parameters

Parameter	Description
<code>input_width</code>	Data input port width
<code>output_width</code>	Data output port width
<code>status_width</code>	Output status port width
<code>config_width</code>	Input configuration port width
<code>max_nfft</code>	The size of the FFT data set is specified as <code>1 << max_nfft</code>
<code>has_nfft</code>	Determines if the size of the FFT can be run time configurable
<code>channels</code>	Number of channels
<code>arch_opt</code>	The implementation architecture.
<code>phase_factor_width</code>	Configure the internal phase factor precision
<code>ordering_opt</code>	The output ordering mode
<code>ovflo</code>	Enable overflow mode
<code>scaling_opt</code>	Define the scaling options
<code>rounding_opt</code>	Define the rounding modes
<code>mem_data</code>	Specify using block or distributed RAM for data memory
<code>mem_phase_factors</code>	Specify using block or distributed RAM for phase factors memory
<code>mem_reorder</code>	Specify using block or distributed RAM for output reorder memory
<code>stages_block_ram</code>	Defines the number of block RAM stages used in the implementation
<code>mem_hybrid</code>	When block RAMs are specified for data, phase factor, or reorder buffer, <code>mem_hybrid</code> specifies where or not to use a hybrid of block and distributed RAMs to reduce block RAM count in certain configurations
<code>complex_mult_type</code>	Defines the types of multiplier to use for complex multiplications
<code>butterfly_type</code>	Defines the implementation used for the FFT butterfly

These parameters are defined in `param_t` in `hls_fft.h`

FFT Configuration Defaults and Valid Values

Parameter	C Type	Default Value	Valid Values
<code>input_width</code>	unsigned	16	8-34
<code>output_width</code>	unsigned	16	input_width to (input_width + max_nfft + 1)
<code>status_width</code>	unsigned	8	Depends on FFT configuration
<code>config_width</code>	unsigned	16	Depends on FFT configuration
<code>max_nfft</code>	unsigned	10	3-16
<code>has_nfft</code>	bool	false	True, False
<code>channels</code>	unsigned	1	1-12
<code>arch_opt</code>	unsigned	pipelined_streaming_io	automatically_select, pipelined_streaming_io, radix_4_burst_io, radix_2_burst_io, radix_2_lite_burst_io
<code>phase_factor_width</code>	unsigned	16	8-34
<code>ordering_opt</code>	unsigned	bit_reversed_order	bit_reversed_order natural_order
<code>ovflo</code>	bool	true	false true
<code>scaling_opt</code>	unsigned	scaled	Scaled, unscaled, block_floating_point
<code>rounding_opt</code>	unsigned	truncation	truncation convergent_rounding
<code>mem_data</code>	unsigned	block_ram	block_ram distributed_ram
<code>mem_phase_factors</code>	unsigned	block_ram	block_ram distributed_ram
<code>mem_reorder</code>	unsigned	block_ram	block_ram distributed_ram
<code>stages_block_ram</code>	unsigned	(max_nfft < 10) ? 0 : (max_nfft - 9)	0-11
<code>mem_hybrid</code>	bool	false	false true
<code>complex_mult_type</code>	unsigned	use_mults_resources	use_luts use_mults_resources use_mults_performance
<code>butterfly_type</code>	unsigned	use_luts	use_luts use_xtremedsp_slices

FFT Function: Data

► Input & Output Data

- Data Input and Output
 - Array arguments of type ap_fixed or float type
 - Two-dimensional arrays are used for multi-channel functionality
 - First dimension represents the channels; Second dimension the data for each channel
 - IP only supports multi-channels for ap_fixed (not float)
- Known Limitation in 2013.3
 - If float types are used the variable must be defined using the static qualifier

```
typedef float data_t;

static complex<data_t> fft_in[FFT_LENGTH];
static complex<data_t> fft_out[FFT_LENGTH];
```

```
#include "hls_fft.h"

hls::fft<STATIC_PARAM> (
    INPUT_DATA_ARRAY,
    OUTPUT_DATA_ARRAY,
    OUTPUT_STATUS,
    INPUT_RUN_TIME_CONFIGURATION);
```

// Static Parameterization Struct
// Input data fixed or float
// Output data fixed or float
// Output Status
// Input Run Time Configuration

FFT Function: Run Time Configuration

► Dynamic Configuration

- Parameter INPUT_RUN_TIME_CONFIGURATION can be changed at run time
- Access is made using the pre-defined struct **hls::ip_fft::config_t<STATIC_PARAM>**
 - Notice, this also uses the FFT parameterization struct

```
#include "hls_fft.h"

struct param1 : hls::ip_fft::params_t {
    static const unsigned ordering_opt = hls::ip_fft::natural_order;
    static const unsigned config_width = FFT_CONFIG_WIDTH;
    static const unsigned status_width = FFT_STATUS_WIDTH;
};

typedef hls::ip_fft::config_t<param1> config_t;
config_t fft_config1;
```

```
#include "hls_fft.h"

hls::fft<STATIC_PARAM> (
    INPUT_DATA_ARRAY,
    OUTPUT_DATA_ARRAY,
    OUTPUT_STATUS,
    INPUT_RUN_TIME_CONFIGURATION); // Input Run Time Configuration
```

```
// Set FFT length to 512 => log2(512) =>9
fft_config1->setNfft(9);

// Forward FFT
fft_config1->setDir(0);
// Inverse FFT
fft_config1->setDir(1);

// Set FFT Scaling
fft_config1->setSch(0x2AB);

// FFT IP
hls::fft<param1> (xn1, xk1, &fft_status1, &fft_config1);
```

FFT Function: Run Time Output Status

► Dynamic Status Monitoring

- Parameter OUTPUT_STATUS can be read at run time
- Access is made using the pre-defined
struct **hls::ip_fft::status_t<STATIC_PARAM>**
 - Notice, this also uses the FFT parameterization struct

```
#include "hls_fft.h"

struct param1 : hls::ip_fft::params_t {
    static const unsigned ordering_opt = hls::ip_fft::natural_order;
    static const unsigned config_width = FFT_CONFIG_WIDTH;
    static const unsigned status_width = FFT_STATUS_WIDTH;
};

typedef hls::ip_fft::status_t<param1> status_t;
status_t fft_status1;
```

```
#include "hls_fft.h"

hls::fft<STATIC_PARAM> (
    INPUT_DATA_ARRAY,
    OUTPUT_DATA_ARRAY,
    OUTPUT_STATUS, // Output Status
    INPUT_RUN_TIME_CONFIGURATION); // Input Run Time Configuration
```

```
// FFT IP
hls::fft<param1> (xn1, xk1, &fft_status1, &fft_config1);

// Check the overflow flag
bool *ovflo = fft_status1->getOvflo();

// Obtain the block exponent
unsigned int *blk_exp = fft_status1->getBlkExp();
```

FIR Function

➤ Using the FIR

- Include the `hls_fir.h` library in the code
 - This defines the FIR and supporting structs and types
 - Allows `hls::FIR` to be instantiated in your code
 - Unlike the FFT, the FIR is instantiated as a class and executed with the `run` method
- Create the `STATIC_PARAM` template parameter to configure the FIR
 - The `STATIC_PARAM` template parameter defines all static configuration values
 - The library provides a pre-defined struct `hls::ip_fir::params_t` to perform this
- There are no default values for the Coefficients
 - You Must Always create a user defined struct based on `hls::ip_fir::params_t`

```
#include "hls_fir.h"  
  
// Create an instance of the FIR  
static hls::FIR<STATIC_PARAM> fir1;           // Static parameterization  
  
// Execute the FIR instance fir1  
fir1.run(INPUT_DATA_ARRAY,  
          OUTPUT_DATA_ARRAY);  
          // Input Data  
          // Output Data
```

FIR Static Configuration

► Default Configuration Pre-defined

- The `hls_fir.h` header file defines struct `hls::ip::fir::params_t`
- This contains the default values for configuring the FIR

► User defined struct must be created

- Allows the FIR to be easily configured
- Based on `hls::ip::fir::params_t`
- User struct simply updates any specified value

```
#include "hls_fir.h"

struct param1 : hls::ip::fir::params_t {
    static const double coeff_vec[sg_fir_srrc_coeffs_len];
    static const unsigned num_coeffs = sg_fir_srrc_coeffs_len;
    static const unsigned input_width = INPUT_WIDTH;
    static const unsigned quantization = hls::ip::fir::quantize_only;};

// Execute FIR
static hls::FIR<param1> fir1;
fir1.run(fir_in, fir_out);
```

Include Header File

Create struct based on `hls::ip::fir::params_t`

- `params_t` includes the default values
- Struct `param1` overrides 3 default value in this example and specifies the coefficients
 - The `coeff_vec` has no default
- Struct `param1` is then used to configure the FIR
- Used as `STATIC_PARAM` in previous slide

FIR Configuration Parameters

Parameter	Description
<code>input_width</code>	Data input port width
<code>input_fractional_bits</code>	Number of fractional bits on the input port
<code>output_width</code>	Data output port width
<code>output_fractional_bits</code>	Number of fractional bits on the output port
<code>coeff_width</code>	Bit-width of the coefficients
<code>coeff_fractional_bits</code>	Number of fractional bits in the coefficients
<code>num_coeffs</code>	Number of coefficients
<code>coeff_sets</code>	Number of coefficient sets
<code>input_length</code>	Number of samples in the input data
<code>output_length</code>	Number of samples in the output data
<code>num_channels</code>	Specify the number of channels of data to process
<code>total_number_coeff</code>	Total number of coefficients
<code>coeff_vec[total_num_coeff]</code>	The coefficient array
<code>filter_type</code>	The type implementation used for the filter
<code>rate_change</code>	Specifies integer or fractional rate changes
<code>interp_rate</code>	The interpolation rate
<code>decim_rate</code>	The decimation rate
<code>zero_pack_factor</code>	Number of zero coefficients used in interpolation
<code>rate_specification</code>	Specify the rate as frequency or period
<code>hardware_oversampling_rate</code>	Specify the rate of over-sampling
<code>sample_period</code>	The hardware oversample period
<code>sample_frequency</code>	The hardware oversample frequency
<code>quantization</code>	The quantitation method to be used
<code>best_precision</code>	Enable or disable the best precision
<code>coeff_structure</code>	The type of coefficient structure to be used
<code>output_rounding_mode</code>	Type of rounding used on the output
<code>filter_arch</code>	Selects a systolic or transposed architecture
<code>optimization_goal</code>	Specify a speed or area goal for optimization
<code>inter_column_pipe_length</code>	The pipeline length required between DSP columns
<code>column_config</code>	Specifies the number of DSP48 column
<code>config_method</code>	Specifies how the DSP48 columns are configured
<code>coeff_padding</code>	Number of zero padding added to the front of the filter

These parameters are defined in `param_t` in `hls_fir.h`

FIR Configuration Defaults and Valid Values

Parameter	C Type	Default Value	Valid Values
input_width	unsigned	16	No limitation
input_fractional_bits	unsigned	0	Limited by size of input_width
output_width	unsigned	24	No Limitation
output_fractional_bits	unsigned	0	Limited by size of output_width
coeff_width	unsigned	16	No Limitation
coeff_fractional_bits	unsigned	0	Limited by size of coeff_width
num_coeffs	bool	21	Full
coeff_sets	unsigned	1	1-1024
input_length	unsigned	21	No Limitation
output_length	unsigned	21	No Limitation
num_channels	unsigned	1	1-1024
total_number_coeff	unsigned	21	num_coeffs * coeff_sets
coeff_vec[total_num_coeff]	double array	None	Not applicable
filter_type	unsigned	single_rate	single_rate, interpolation, decimation, hilbert, interpolated
rate_change	unsigned	integer	integer, fixed_fractional
interp_rate	unsigned	1	2-1024
decim_rate	unsigned	1	2-1024
zero_pack_factor	unsigned	1	2-8
rate_specification	unsigned	period	frequency, period
hardware_oversampling_rate	unsigned	1	No Limitation
sample_period	bool	1	No Limitation
sample_frequency	unsigned	0.001	No Limitation
quantization	unsigned	integer_coefficients	integer_coefficients, quantize_only, maximize_dynamic_range
best_precision	unsigned	false	false true
coeff_structure	unsigned	non_symmetric	inferred, non_symmetric, symmetric, negative_symmetric, half_band, hilbert
output_rounding_mode	unsigned	full_precision	full_precision, truncate_lsb, non_symmetric_rounding_down, non_symmetric_rounding_up, symmetric_rounding_to_zero, symmetric_rounding_to_infinity, convergent_rounding_to_even, convergent_rounding_to_odd
filter_arch	unsigned	systolic_multiply_accumulate	systolic_multiply_accumulate, transpose_multiply_accumulate
optimization_goal	unsigned	area	area, speed
inter_column_pipe_length	unsigned	4	1-16
column_config	unsigned	1	Limited by number of DSP48s used
config_method	unsigned	single	single, by_channel
coeff_padding	bool	false	false true

FIR Function: Data

➤ Input & Output Data

- Data Input and Output
 - Array arguments

➤ Multi-channel functionality

- Supported through interleaving the data

Interleave 2 channels into the FIR input

```
for (unsigned i = 0; i < LENGTH; ++i) {  
    fir_input_data[2*i] = din_ch1[i];  
    fir_input_data[2*i + 1] = din_ch2[i];  
}
```

```
#include "hls_fir.h"  
  
// Create an instance of the FIR  
static hls::FIR<STATIC_PARAM> fir1;           // Static parameterization  
  
// Execute the FIR instance fir1  
fir1.run(INPUT_DATA_ARRAY,                      // Input Data  
          OUTPUT_DATA_ARRAY);                   // Output Data
```

De-Interleave 2 channels from the FIR output

```
for(unsigned i = 0; i < LENGTH; ++i) {  
    dout_ch1[i] = fir_output_data[2*i];  
    dout_ch2[i] = fir_output_data[2*i+1];  
}
```

FIR Function: Run Time Configuration

► Dynamic Configuration

- Parameter INPUT_RUN_TIME_CONFIGURATION can be changed at run time
- In some modes a dynamic input can control how the coefficients are used
 - This is a standard 8-bit input

```
#include "hls_fir.h"

struct param1 : hls::ip_fir::params_t {
    static const double coeff_vec[sg_fir_srrc_coeffs_len];
    static const unsigned num_coeffs = sg_fir_srrc_coeffs_len;
    static const unsigned input_width = INPUT_WIDTH;
    static const unsigned quantization = hls::ip_fir::quantize_only;
};
```

```
#include "hls_fir.h"

// Create an instance of the FIR
static hls::FIR<STATIC_PARAM> fir1; // Static parameterization

// Execute the FIR instance fir1
fir1.run(INPUT_DATA_ARRAY, // Input Data
         OUTPUT_DATA_ARRAY, // Output Data
         INPUT_RUN_TIME_CONFIGURATION); // Run time input configuration
```

```
// Define the configuration type
typedef ap_uint<8> config_t;

// Define the configuration variable
config_t fir_config = 8;

// Use the configuration in the FFT
static hls::FIR<param1> fir1;
fir1.run(fir_in, fir_out, &fir_config);
```

Using the FFT and FIR IP

➤ FFT and FIR support pipelined implementations

- The functions themselves cannot be pipelined
- They should be parameterized for pipelined operation

➤ The data arguments are always arrays

- These will be implemented as AXI4 Streams in the RTL
 - By default, arrays are implemented as BRAM interfaces

➤ Recommendation

- Use these IP in regions where dataflow optimization is used
- This will auto-convert the input and output arrays into streaming arrays

➤ Alternatively, a Requirement:

- The input and output arrays must be marked as streaming using the command `set_directive_stream` (pragma STREAM)

Outline

- Language Support
- Pointers
- Coding Considerations and IO
- Streams and Shift Registers
- *Libraries Support*
 - FFT & FIR
 - OpenCV
- Summary

Video Libraries Supported in Vivado HLS

➤ C video libraries

- Available within Vivado HLS tool header files
 - *hls_video.h* library
 - *hls_opencv.h* library

➤ Enable migration of OpenCV designs for use with the Vivado HLS tool

➤ HLS video library is intended to replace many basic OpenCV functions

- Similar interfaces and algorithms to OpenCV
- Focus on image processing functions implemented in FPGA fabric such as full HD video processing
- Includes FPGA-specific optimizations
 - Fixed-point operations instead of floating point
 - On-chip line buffers and window buffers
- Not necessarily bit-accurate
- Libraries support standard AXI4 interfaces for easy system integration

Video Library Functions

- C++ code contained in hls namespace: `#include "hls_video.h"`
- Similar interface; equivalent behavior with OpenCV
 - OpenCV library: `cvScale(src, dst, scale, shift);`
 - HLS video library: `hls::Scale<...>(src, dst, scale, shift);`
- Some constructor arguments have corresponding or replacement template parameters
 - OpenCV library: `cv::Mat mat(rows, cols, CV_8UC3);`
 - HLS video library: `hls::Mat<ROWS, COLS, HLS_8UC3> mat(rows, cols);`
 - ROWS and COLS specify the maximum size of an image processed

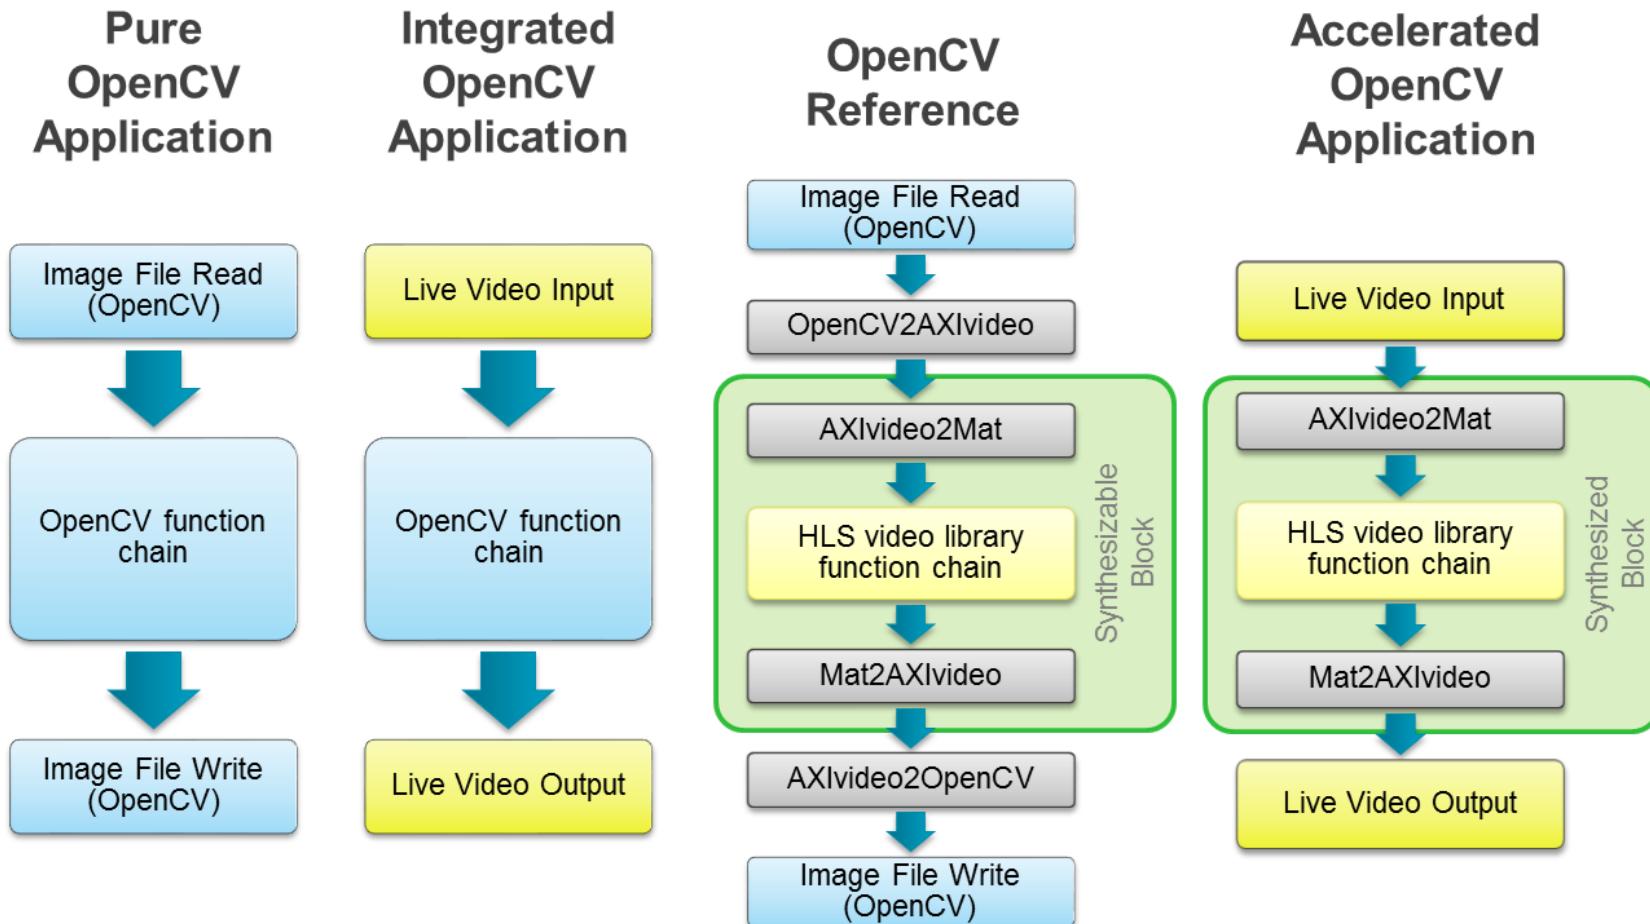
Video Library Supported Functions

Video Data Modeling		AXI4-Stream IO Functions	
Linebuffer class	Window class	AXIvideo2Mat	Mat2AXIvideo

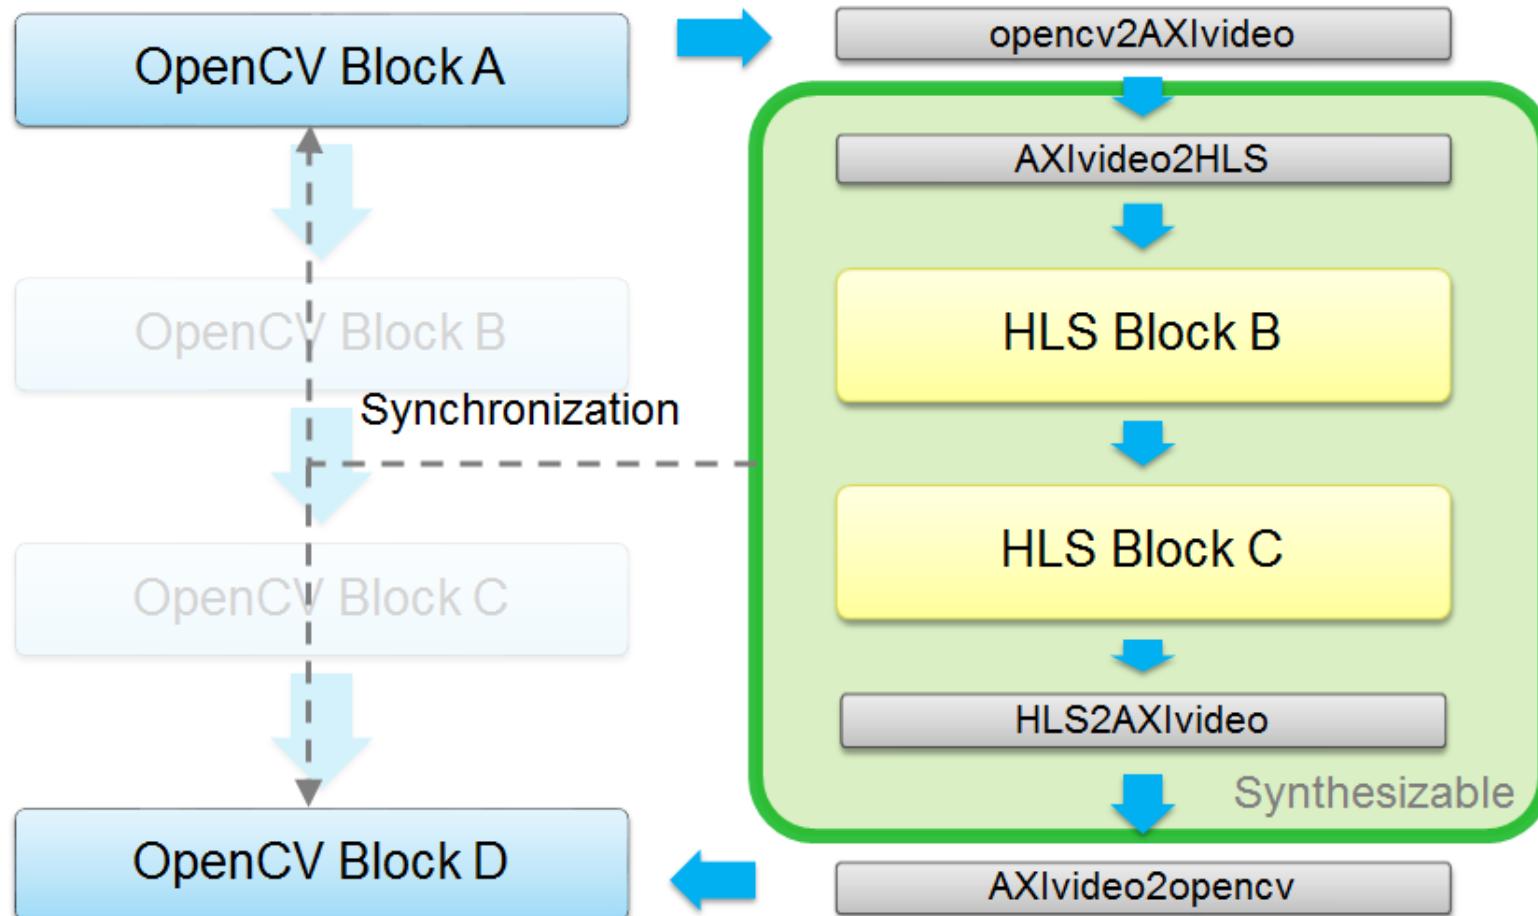
OpenCV Interface Functions			
cvMat2AXIvideo	AXIvideo2cvMat	cvMat2hlsMat	hlsMat2cvMat
IplImage2AXIvideo	AXIvideo2IplImage	IplImage2hlsMat	hlsMat2IplImage
CvMat2AXIvideo	AXIvideo2CvMat	CvMat2hlsMat	hlsMat2CvMat

Video Functions			
AbsDiff	Duplicate	MaxS	Remap
AddS	EqualizeHist	Mean	Resize
AddWeighted	Erode	Merge	Scale
And	FASTX	Min	Set
Avg	Filter2D	MinMaxLoc	Sobel
AvgSdv	GaussianBlur	MinS	Split
Cmp	Harris	Mul	SubRS
CmpS	HoughLines2	Not	SubS
CornerHarris	Integral	PaintMask	Sum
CvtColor	InitUndistortRectifyMap	Range	Threshold
Dilate	Max	Reduce	Zero

Using OpenCV in FPGA Designs

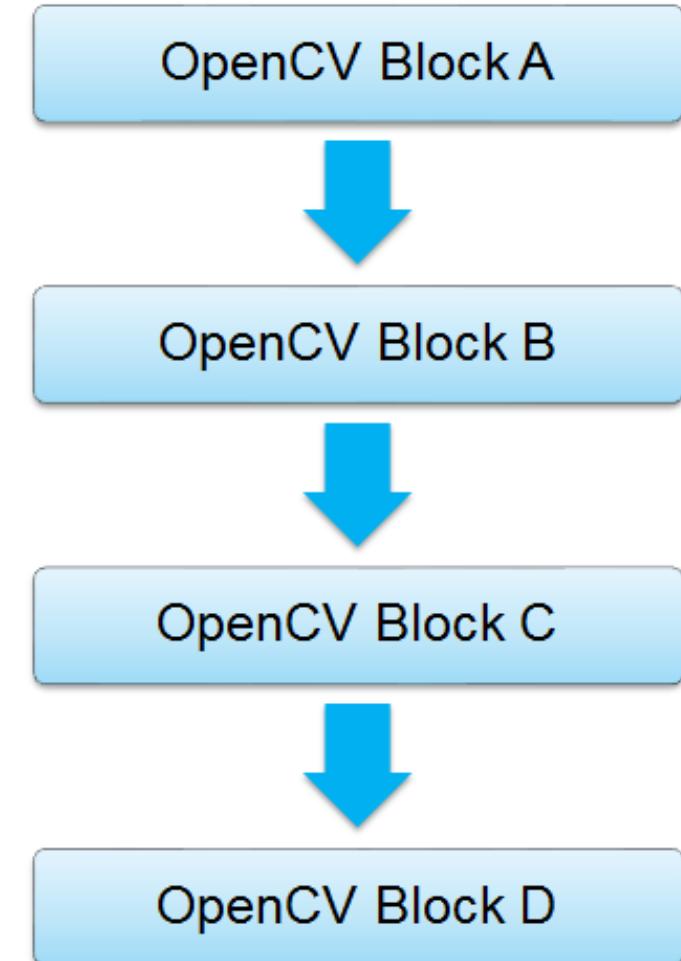


Partitioned OpenCV Application



OpenCV Design Flow

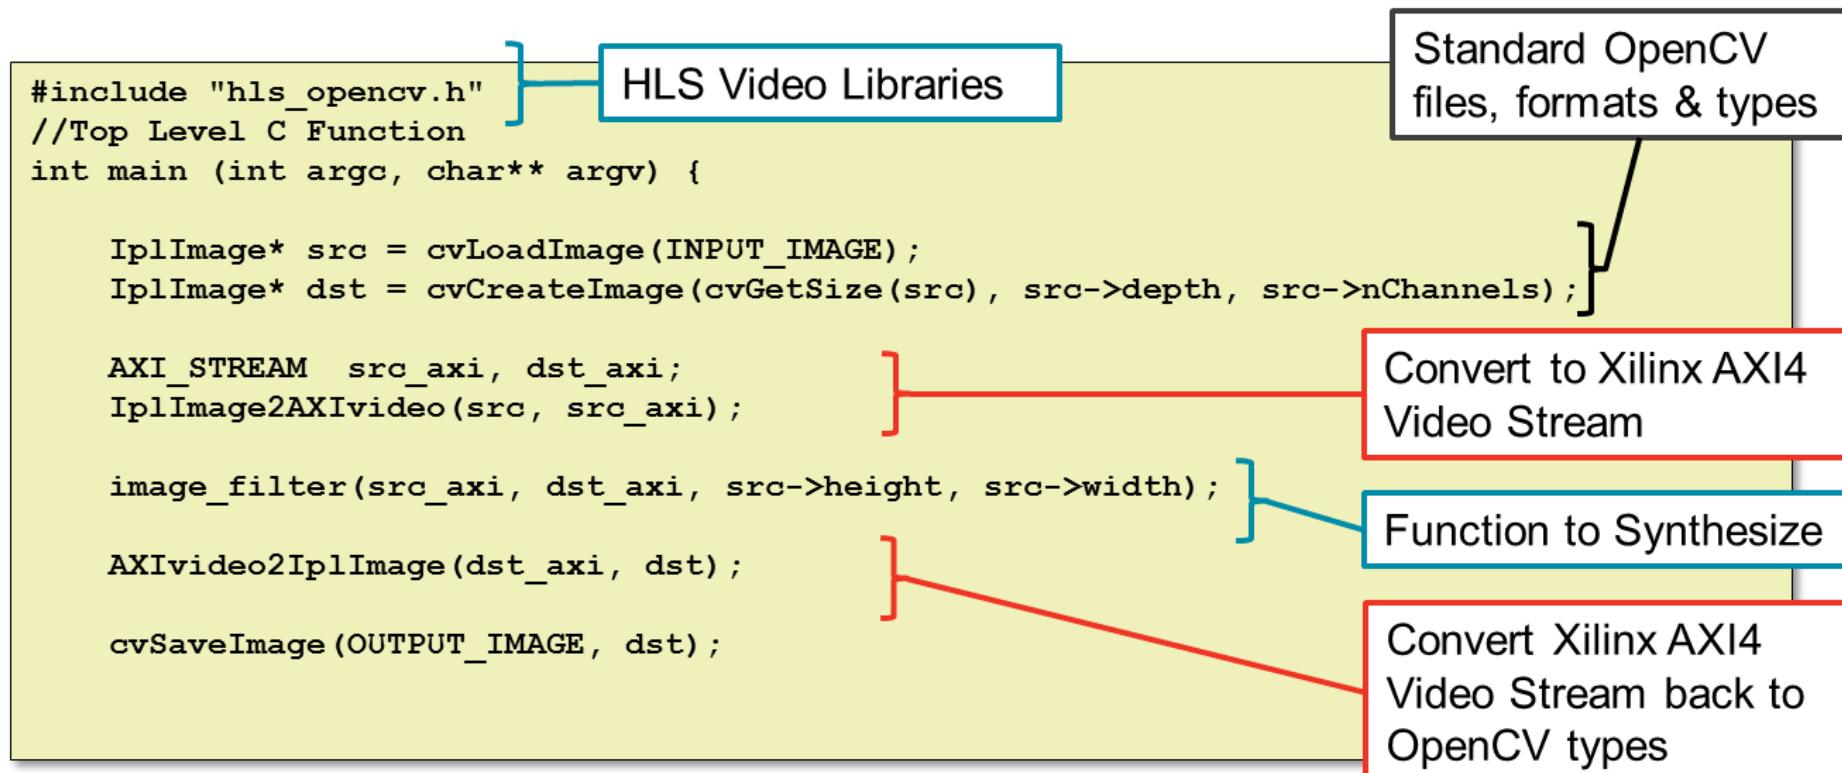
- Develop OpenCV application on desktop
- Run OpenCV application on ARM® cores without modification
- Abstract FPGA portion using I/O functions
- Replace OpenCV function calls with synthesizable code
- Run HLS to generate FPGA accelerator
- Replace call to synthesizable code with call to FPGA accelerator



C Testbench: Interface Library Example

► Interface libraries convert to/from OpenCV image to HLS type

- HLS MAT format: synthesizable and AXI4 Stream support



Synthesizable C Function Example

```
#include "hls_video.h"
#include "ap_axi_sdata.h";
//Top Level C Function for Synthesis
void image_filter(AXI_STREAM& inter_pix, AXI_STREAM& out_pix, int rows, int cols) {
    //Create AXI streaming interfaces for the core

    RGB_IMAGE img_0(rows, cols);
    ..etc..
    RGB_IMAGE img_5(rows, cols);
    RGB_PIXEL pix(50, 50, 50);
#pragma HLS dataflow
    hls::AXIvideo2Mat(inter_pix, img_0);

    hls::Sobel(img_0, img_1, 1, 0);
    hls::SubS(img_1, pix, img_2);
    hls::Scale(img_2, img_3, 2, 0);
    hls::Erode(img_3, img_4);
    hls::Dilate(img_4, img_5);

    hls::Mat2AXIvideo(img_5, out_pix);
}
```

HLS Video & AXI Struct Libraries

Convert Xilinx AXI4 Video Stream to HLS Mat data type

HLS Video functions are drop-in replacement for OpenCV function & provide high QoR

Convert HLS Mat type to Xilinx AXI4 Video Stream

Outline

- Language Support
- Pointers
- Coding Considerations and IO
- Streams and Shift Registers
- Libraries Support
 - FFT & FIR
 - OpenCV
- *Summary*

Summary

➤ Vivado HLS supports C, C++, and SystemC

- Vast majority of constructs are synthesizable provided they are statically defined at compile time

➤ Some of the unsupported constructs are

- System calls and function pointer
- Forward declared data type
- Recursive type functions
- Dynamic memory allocation
- Dynamic virtual calls
- Pointer reinterpretation

Summary

- **Multi-access pointer is a pointer which is read and written to multiple times**
 - Must be marked as volatile or reads and writes will be optimized away
 - Cannot be rigorously validated with C simulation
 - Require a depth specification for RTL simulation
- **When struct is passed as a pointer the data are accessed as pass-by-reference**
- **When struct is not passed a pointer the data are accessed as by value**
- **When pointer arithmetic is performed, argument may be implemented as ap_bus or ap_fifo**
- **Stream can be used instead of multiple-access pointer**
- **Stream and shift registers classes may model and implement hardware better**
- **HLS video libraries**
 - Drop-in replacement for OpenCV and provide high QoR